

COMP 6321 Machine Learning

Kernel Density Estimation

Computer Science & Software Engineering
Concordia University, Fall 2020



Parametric vs non-parametric

Parametric models have a fixed number of adaptable parameters, independent of the amount of data.

- Logistic regression $\mathbf{w} \in \mathbb{R}^D$
- K -component mixture of Gaussians $\{\boldsymbol{\pi}_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k\}_{k=1}^K$


Non-parametric models have a variable number of parameters, adapting to the amount of data.

- Kernel density estimators (this lecture)
- Support vector machines (next lecture)
- Nearest neighbour classifiers (next lecture after that)

Aside: In practical terms, the distinction is somewhat fuzzy. Example: if a model selection procedure tends to select larger parametric models given more data, are we ‘adapting’ the number of parameters to the data and therefore “less parametric”? Histograms are considered “non-parametric” because they make fewer assumptions, despite having a fixed number of parameters (bin frequencies).




Kernel Density Estimation



Density Estimation


- **Goal:** Given samples $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ we want to estimate $p(\mathbf{x})$ at a new point \mathbf{x} .
 - for scoring examples, or for novelty/anomaly detection...
 - for classification by comparing $p(\mathbf{x})$ to some $p(\mathbf{y})$
- Example: Fit a GMM to \mathcal{D} and evaluate $p(\mathbf{x} \mid \boldsymbol{\pi}, \boldsymbol{\mu}, \boldsymbol{\Sigma})$
- Example: Fit a histogram to \mathcal{D} and evaluate density of bin that \mathbf{x} falls into.

An illustration of the histogram approach to density estimation, in which a data set of 50 data points is generated from the distribution shown by the green curve. Histogram density estimates are shown for various values of bin width Δ




Kernel Density Estimation


- **Idea:** Estimate $p(\mathbf{x})$ by “smoothing” the data \mathcal{D} itself. Do this by convolving the \mathbf{x}_i with some “kernel.”
 - The model of $p(\mathbf{x})$ is expressed directly in terms of data \mathbf{x}_i
 - The complexity of the density estimate can scale with the amount of data (more data => more terms in $p(\mathbf{x})$)



GMM densities




Synthetic data from GMM



Density from GMM fit to the synthetic data


$p(\mathbf{x})$
(really $p(\mathbf{x} \mid \boldsymbol{\pi}, \boldsymbol{\mu}, \boldsymbol{\Sigma})$)

$\boldsymbol{\pi}_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k$




Model parameters of the GMM that was fit


GMM densities (parametric)



$K=1$ component




$K=3$ components




$K=20$ components

Histogram densities


(non-parametric,
but not kernel
density)



binwidth=3.0





binwidth=1.0




binwidth=0.3


Kernel densities ("tophat")



$h=3.0$ bandwidth





$h=1.0$ bandwidth




$h=0.3$ bandwidth


Kernel densities ("gaussian")



$h=3.0$ bandwidth



$h=1.0$ bandwidth



$h=0.3$ bandwidth

Kernel density estimator


A **kernel density estimator** (or **Parzen estimator**) is computed by adding an instance of the kernel function *centered at each data point* \mathbf{x}_i and then normalizing.

$$p(\mathbf{x}) = \frac{1}{N} \sum_{i=1}^N k(\mathbf{x} - \mathbf{x}_i)$$

A valid “kernel function” must satisfy:

$$k(\mathbf{x}) \geq 0 \quad (\text{non-negative})$$


$$\int_{\mathbb{R}^D} k(\mathbf{x}) d\mathbf{x} = 1 \quad (\text{normalized})$$




With Gaussian kernel, it's just an N -component GMM with equal weights and variances! “non-parametric”

Kernels for density estimation

Gaussian kernel $k(\mathbf{x}) = \frac{1}{(2\pi)^{\frac{D}{2}}} \exp\left(-\frac{1}{2} \|\mathbf{x}\|^2\right)$ i.e. $\mathcal{N}(\mathbf{x} \mid \mathbf{0}, \mathbf{I})$




Tophat kernel $k(\mathbf{x}) = \begin{cases} \frac{\Gamma(\frac{D}{2}+1)}{\sqrt{\pi}^D} & \|\mathbf{x}\|^2 \leq 1 \\ 0 & \text{otherwise} \end{cases}$



where Γ is the gamma function.
e.g. when $D = 1$

$$\frac{\Gamma(\frac{3}{2})}{\sqrt{\pi}} = \frac{1}{2}$$

Parzen window $k(\mathbf{x}) = \begin{cases} 1 & |x_i| \leq \frac{1}{2} \text{ for } i = 1, \dots, D \\ 0 & \text{otherwise} \end{cases}$






We can “add bandwidth h ” to any kernel $k(\mathbf{x})$ by taking $k_h(\mathbf{x}) = \frac{1}{h^D} k\left(\frac{\mathbf{x}}{h}\right)$

Smoothing as convolution

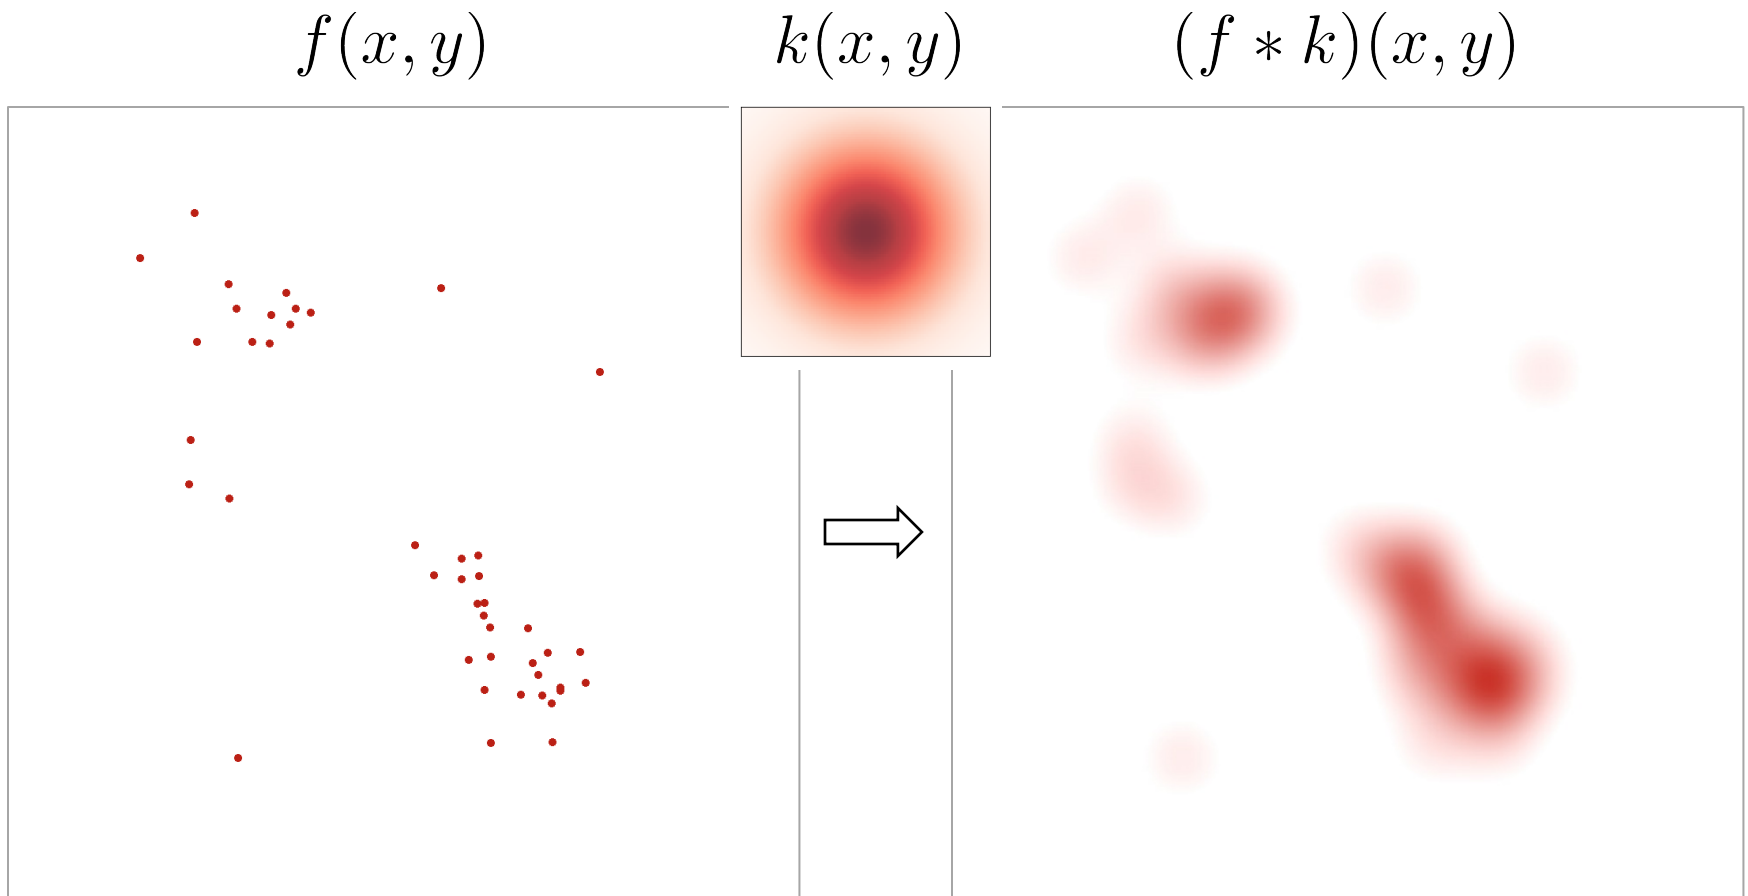
$$f(x, y)$$

$$k(x, y)$$

$$(f * k)(x, y)$$



Smoothing as convolution



" f is zero everywhere except at the data points"

Kernel density as convolution

The **empirical density** is zero everywhere except the at the empirically observed data:

$$f(\mathbf{x}) = \frac{1}{N} \sum_{i=1}^N \delta(\mathbf{x} - \mathbf{x}_i)$$

The **kernel density** is the empirical density convolved with a smoothing kernel k :


$$\begin{aligned} p(\mathbf{x}) &= (f * k)(\mathbf{x}) \quad \text{“convolution of } f \text{ by } k\text{”} \\ &= \frac{1}{N} \sum_{i=1}^N k(\mathbf{x} - \mathbf{x}_i) \end{aligned}$$

(proof of this on blackboard if time)

The “Dirac delta function” satisfies:

$$\delta(\mathbf{x}) = \begin{cases} +\infty & \mathbf{x} = \mathbf{0} \\ 0 & \mathbf{x} \neq \mathbf{0} \end{cases}$$

$$\int_{\mathbb{R}^D} \delta(\mathbf{x}) d\mathbf{x} = 1$$




Kernel Parameters



Two free parameters in this “non-parametric” model

1. Kernel function:

These work well and have nice probabilistic interpretations, but may be slow to estimate predict new $p(\mathbf{x})$ because *all* training points contribute to density, even if very small contribution.



These are faster at predicting a new density $p(\mathbf{x})$ because the kernels have a limited range of support so only a small subset of training points will contribute to density.

2. Kernel function bandwidth h

Scikit-learn and Numpy

```
gmm = sklearn.mixture.GaussianMixture(n_components=k)
gmm.fit(X)
```


(included for completeness – not a kernel density estimator!)

```
ranges = [[-10, 10], [-10, 10]]
bins = np.ptp(ranges, axis=1)/binwidth
hist = np.histogram2d(X[:,0], X[:,1], bins=bins, range=ranges, density=1)[0]
```

(included for completeness – not a kernel density estimator!)

```
kde = sklearn.neighbors.KernelDensity(bandwidth=bandwidth)
kde.fit(X)
```

(use kernel='tophat' etc to choose non-Gaussian kernel)



PRML Readings

§2.5.0 Nonparametric methods

§2.5.1 Kernel density estimators